

Module 7: Memory Management

Reading: Chapter 8

- **To provide a detailed description of various ways of organizing memory hardware.**
- **To discuss various memory-management techniques, including paging and segmentation.**

Virtual memory	Physical memory
0	0
1	1
2	2
3	3
4	4
5	5
6	6
7	7
8	8
9	9
10	10
11	11
12	12
13	13
14	14
15	15
16	16

Memory

Management

Paging

Segmentation

Continuous
Memory Allocation

Background

Basic
Hardware

Address
Binding

Physical vs
Logical
Addressing

Swapping

Memory Management

Challenge:

- **To accommodate multiple processes in main memory...**
 - Allocate and deallocate the memory for them
 - Make sure they do not corrupt each other's memory
 - Make sure that all of this is done efficiently and transparently
 - Not concerned with how addresses are generated.
- **Main requirements**
 - Relocation.
 - Sharing.
 - Allocation.
 - Protection

Swapping

Run out of memory?

- Move process out of main memory


Where?

- **Backing store** – fast disk large enough to accommodate copies of all memory images for all users; must provide direct access to these memory images

Modified versions of swapping are found on many systems (i.e., UNIX, Linux, and Windows)

- Major part of swap time is transfer time; total transfer time is directly proportional to the amount of memory swapped
- Needs intelligent decisions, if serious swapping starts, you can go for a coffee

Schematic View of Swapping



RAM memory	RAM memory
0	0
1	1
2	2
3	3
4	4
5	5
6	6
7	7
8	8
9	9
10	10
11	11
12	12
13	13
14	14

Memory

Management

Paging

Segmentation

Continuous
Memory Allocation

Mapping
and
Protection

Memory
Allocation

Fragmentation

Background

Basic
Hardware

Address
Binding

Physical vs
Logical
Addressing

Swapping

Contiguous Allocation

What do we mean by contiguous allocation?

- Each process occupies a contiguous block of physical memory

Where is the OS kernel?


- usually held in low memory, together with the interrupt vector

With contiguous allocation, a base and a limit register are sufficient to describe the address space of a process


Fixed Partitioning or Dynamic Partitioning can be used

Memory Allocation: Fixed Partitioning

- Partition main memory into a set of non overlapping regions called **partitions**
- Partitions can be of equal or unequal sizes
- any process whose size is less than or equal to a partition size can be loaded into the partition
- if all partitions are occupied, the operating system can swap a process out of a partition



Equal-size partitions





Unequal-size partitions

Placement Algorithm with Partitions

- **Equal-size partitions**
 - If there is an available partition, a process can be loaded into that partition
 - because all partitions are of equal size, it does not matter which partition is used
 - If all partitions are occupied by blocked processes, choose one process to swap out to make room for the new process

Placement Algorithm with Partitions – unequal sizes




Fixed Partitioning – Problems

- Main memory use is inefficient. Any program, no matter how small, occupies an entire partition. This is called **internal fragmentation**.
- Unequal-size partitions lessens these problems but they still remain...
- Equal-size partitions was used in early IBM's OS/MFT (Multiprogramming with a Fixed number of Tasks)


Memory Allocation: Dynamic Allocation

Idea: Forget partitions, put a new process into a large enough hole of unused memory.

- The physical memory consists of allocated partitions and *holes* – blocks of available memory between already allocated space.
- Initially, there is one huge hole containing all free memory
- Operating system maintains information about:
a) allocated partitions b) free partitions (holes)




Dynamic Allocation: an example



- A hole of 64K is left after loading 3 processes: not enough room for another process
- Eventually each process is blocked. The OS swaps out process 2 to bring in process 4

Dynamic Allocation: an example




- another hole of 96K is created
- Eventually each process is blocked. The OS swaps out process 1 to bring in again process 2 and another hole of 96K is created...

Placement Algorithm in Dynamic Allocation

How to satisfy a request of size n from a list of free holes?

- **First-fit:** Allocate the *first* hole that is big enough
- **Next-fit:** Allocate the first hole that is big enough, starting from the last allocated partition
- **Best-fit:** Allocate the *smallest* hole that is big enough; must search entire list, unless ordered by size. Produces the smallest leftover hole.
- **Worst-fit:** Allocate the *largest* hole; must also search entire list. Produces the largest leftover hole.

First-fit and best-fit better than worst-fit and next-fit in terms of speed and storage utilization



Example Memory Configuration Before and After Allocation of 16 Kbyte Block

Fragmentation

What happens after some time of allocation and deallocation?

- **Plenty of small holes nobody can fit into**
- **External Fragmentation** – total memory space exists to satisfy a request, but it is not contiguous
- **Reduce external fragmentation by compaction**
 - **Shuffle memory contents to place all free memory together in one large block**
 - **Compaction is possible *only* if relocation is dynamic, and address binding is done at execution time**
 - **I/O problem**
 - **Latch job in memory while it is involved in I/O**
 - **Do I/O only into OS buffers**

What happens if we allocate more than requested?

- **e.g. we allocate only in 1k blocks, or add the remaining small hole to the process's partition**
- **Internal Fragmentation**

Memory Management

Paging

Segmentation

Continuous
Memory Allocation

Mapping
and
Protection

Memory
Allocation

Fragmentation

Program
Segments

Segment
Table

Segmentation
Hardware

Examples

Background

Basic
Hardware

Address
Binding

Physical vs
Logical
Addressing

Swapping

Non-contiguous Allocation

OK, contiguous allocation seems natural, but we have the fragmentation problem.

Also: how do we grow/shrink the address space of a process?

Idea: If we can split the process into several non-contiguous chunks, we would not have these problems


Two principal approaches:

- **Segmentation**
- **Paging**
- **Can be combined**


Segmentation

- **Memory-management scheme that supports user view of memory**
- **A program is a collection of segments. A segment is a logical unit such as:**
 - main program,**
 - procedure,**
 - function,**
 - method,**
 - object,**
 - local variables, global variables,**
 - common block,**
 - stack,**
 - symbol table, arrays**


User's View of a Program



Logical View of Segmentation



user space



physical memory space

Segmentation Architecture


- How do we logically address a memory location?
 - Segment number + offset within segment:
 <segment-number, offset>
- What information do we need to store in order to be able to translate logical address into physical?
 - For each segment:
 - Start of the segment (*base*)
 - Length of the segment (*limit*)
- Where do we store this information?
 - Segment table
- What do we need to know about segment table?
 - Where it starts? *Segment-table base register (STBR)*
 - How large is it? *Segment-table length register (STLR)*
 - segment number s is legal if $s < \text{STLR}$

Segmentation Architecture (Cont.)


- **Relocation.**
 - dynamic
 - change the base pointer in the segment table entry
- **Sharing.**
 - shared segments
 - same segment table entries
- **Allocation.**
 - first fit/best fit
 - external fragmentation
- **Protection**
 - protection bits associated with each segment
 - code sharing occurs at segment level

Segmentation Hardware


How to calculate physical address with segmentation?




Details on address translation (in hardware).



Example of Segmentation



Sharing of Segments



Segmentation and protection

- **Each entry in the segment table contains protection information**
 - **Segment length**
 - **User privileges on the segment: read, write, execution.**
 - If at the moment of address translation, it is found that the user does not have access permissions ☐ interruption
 - This info can vary from user to user (i.e. process to process), for the same segment.

limit	base	read, write, execute?
-------	------	-----------------------

Discussion of simple segmentation

- **Advantages: the unit of allocation is**
 - Smaller than the entire program/process
 - A logical unit known by the programmer
 - Segments can change place in memory
 - Easy to protect and share the segments (in principle).
- **Disadvantage: the problems of dynamic partitioning:**
 - External fragmentation has not been eliminated:
 - Holes in memory, compression?
- **Another solution is to try to simplify the mechanism by using small units of allocation of equal size:**
 - **PAGING**

Segmentation versus pagination

- **The problem with segmentation is that the unit of allocation of memory (the segment) is of variable length).**
- **Paging uses units of memory allocation of fixed size, and eliminates this problem.**





Mem. address	Mem. content	Mem. address	Mem. content
0	0.0	0	0.0
1	0.1	1	0.1
2	0.2	2	0.2
3	0.3	3	0.3
4	0.4	4	0.4
5	0.5	5	0.5
6	0.6	6	0.6
7	0.7	7	0.7
8	0.8	8	0.8
9	0.9	9	0.9
10	1.0	10	1.0
11	1.1	11	1.1
12	1.2	12	1.2
13	1.3	13	1.3
14	1.4	14	1.4

Memory Management

Memory Management

Segmentation

Paging



Paging

Paging

Page Tables

Paging Hardware

Page Table Structures

Hierarchical

Inverted

Hashed

Segmentation with paging

- Divide physical memory into fixed-sized blocks called frames (size is power of 2, usually between 512 and 16M)
- Divide logical memory into blocks of the same size called pages
- To run a program of size n pages, we need to find n free frames and assign them to the program
 - Keep track of all free frames
 - Set up a page table to translate logical to physical addresses
- What if the program's size is not exact multiple of a page size?
 - Internal fragmentation
 - On average 50% of a page size per process
 - don't want large page size

Example of process loading

Frame number	Main memory
0	
1	
2	
3	
4	
5	
6	
7	
8	
9	
10	
11	
12	
13	
14	

(a) Fifteen Available Pages

Frame number	Main memory
0	A.0
1	A.1
2	A.2
3	A.3
4	
5	
6	
7	
8	
9	
10	
11	
12	
13	
14	

(b) Load Process A

Frame number	Main memory
0	A.0
1	A.1
2	A.2
3	A.3
4	B.0
5	B.1
6	B.2
7	
8	
9	
10	
11	
12	
13	
14	

(b) Load Process B

Frame number	Main memory
0	A.0
1	A.1
2	A.2
3	A.3
4	B.0
5	B.1
6	B.2
7	C.0
8	C.1
9	C.2
10	C.3
11	
12	
13	
14	

(d) Load Process C

- Now suppose that process B is swapped out

Example of process loading (cont.)

- When process A and C are blocked, the pager loads a new process D consisting of 5 pages
- Process D does not occupy a contiguous portion of memory
- There is no external fragmentation
- Internal fragmentation consist only of the last page of each process

Main memory	
0	A.0
1	A.1
2	A.2
3	A.3
4	
5	
6	
7	C.0
8	C.1
9	C.2
10	C.3
11	
12	
13	
14	

(e) Swap out B

Main memory	
0	A.0
1	A.1
2	A.2
3	A.3
4	D.0
5	D.1
6	D.2
7	C.0
8	C.1
9	C.2
10	C.3
11	D.3
12	D.4
13	
14	

(f) Load Process D

Page Tables

0	0
1	1
2	2
3	3

Process A
page table

0	—
1	—
2	—

Process B
page table

0	7
1	8
2	9
3	10

Process C
page table

0	4
1	5
2	6
3	11
4	12

Process D
page table

13
14


Free frame
list

- The OS now needs to maintain a **page table** for each process
- Each entry of a page table consist of the frame number where the corresponding page is physically located
- The page table is indexed by the page number to obtain the frame number
- A free frame list, available for pages, is maintained


Address Translation Scheme

- **Address (logical) generated by CPU is divided into:**
 - *Page number (p)* – used as an index into a *page table* which contains number of each frame in physical memory
 - *Page offset (d)* – combined with frame number, defines the physical memory address that is sent over the address bus

Address Translation in Paging




Logical-to-Physical Address Translation in Paging




Logical address =
Page# = 1, Offset = 478

000001|0111011110




Note that the physical address is found using a concatenation operation (remember that with segmentation, an add operation is necessary)

Paging Example




Free Frames



Implementation of Page Table

Where are page tables stored?

- Ideally – special fast registers on the CPU
 - But page tables can be really huge
- In reality – in main memory
 - *Page-table base register (PTBR)* points to the page table
 - *Page-table length register (PRLR)* indicates the page table size
- What happens on each memory reference?
 - we have to read page table to figure out address translation
 - every data/instruction access requires two memory accesses. One for the page table and one for the data/instruction.



Implementation of Page Table

Two memory accesses per each memory reference is too much!

What to do?


Idea: Cache page table entries

- **special fast-lookup hardware cache called translation look-aside buffer (TLB)**

How does TLB work?

- **associative memory – checks all entries in parallel whether they contain a pair (page #, frame#) for the needed page#**
- **If hit, we got the frame# and saved one memory access**
- **If miss, still need to go the memory to look-up the page table**

Paging Hardware With TLB




Effective Access Time

- Associative Lookup = ε time units
- Assume memory cycle time is 1 microsecond
- Hit ratio – percentage of times that a page number is found in the associative registers; ratio related to number of associative registers
- Hit ratio = α
- Effective Access Time (EAT)
$$\text{EAT} = (1 + \varepsilon) \alpha + (2 + \varepsilon)(1 - \alpha)$$
$$= 2 + \varepsilon - \alpha$$

Memory Protection

- **We can associate several bits with each page**
 - read only/read-write
 - valid-invalid
 - these bits are stored for each entry of the page table
- **Checking for accessing only process's memory**
 - If page table size is fixed so that it stores the whole process's addressable memory
 - valid-invalid bits
 - If page table size corresponds to the memory the process is really using
 - Compare page# with *page-table length register (PTLR)*


Valid (v) or Invalid (i) Bit In A Page Table



Shared Pages

- **Shared code**
 - **One copy of read-only (reentrant) code shared among processes (i.e., text editors, compilers, window systems).**
- **Private code and data**
 - **Each process keeps a separate copy of private code and data**

Shared Pages Example




Page Table Structure

How big can a page table be?

- 32-bit address space, 4kb page $\rightarrow 2^{20}$ pages = 1M page entries!
- Better come-up with something smart!

- Hierarchical Paging
- Hashed Page Tables
- Inverted Page Tables



Hierarchical Page Tables

Idea: The page table itself is paged

- A simple technique is a two-level page table:
 - The outer page table entries point to the pages of the real page table
- Sometimes (i.e. SPARC), three-level page tables are used


Two-Level Paging Example

- A logical address (on 32-bit machine with 4K page size) is divided into:
 - a page number consisting of 20 bits
 - a page offset consisting of 12 bits
- Since the page table is paged, the page number is further divided into:
 - a 10-bit page number
 - a 10-bit page offset
- Thus, a logical address is as follows:

page number		page offset
p_1	p_2	d
10	10	12


where p_1 is an index into the outer page table, and p_2 is the displacement within the page of the outer page table

Two-Level Page-Table Scheme



Address-Translation Scheme

- Address-translation scheme for a two-level 32-bit paging architecture



- In the case of the two-level paging, the use of the TLB becomes even more important to reduce the multiple memory accesses to calculate the physical address.


Hashed Page Tables

Assume the address space is 64 bits

How many level of page hierarchy would we need?

- **With 4k pages we have 2^{52} pages**
 - **Assuming each entry fits into 4 bytes, 5 levels are still not enough! (one page fits only 1k entries)**
 - **That is not feasible, what can we do?**
-
- **The virtual page number is hashed into a page table. This page table contains a chain of elements hashing to the same location.**
 - **Virtual page numbers are compared in this chain searching for a match. If a match is found, the corresponding physical frame is extracted.**

Hashed Page Table



Inverted Page Table

We can have plenty of processes...

... and each one of them has big page table.

Of course, most of the page table entries are invalid, as most of the processes do not fit into memory

Still, we are wasting memory on their page tables...

Idea: have only one entry for each physical frame of memory

- **Entry consists of the virtual address of the page stored in that real memory location, with information about the process that owns that page**

Main memory	
0	A.0
1	A.1
2	A.2
3	A.3
4	
5	
6	
7	C.0
8	C.1
9	C.2
10	C.3
11	
12	
13	
14	

(e) Swap out B

Main memory	
0	A.0
1	A.1
2	A.2
3	A.3
4	D.0
5	D.1
6	D.2
7	C.0
8	C.1
9	C.2
10	C.3
11	D.3
12	D.4
13	
14	

(f) Load Process D


Inverted Page Table

**Now we need only one entry per physical frame,
not per each process's logical page**

But how do we translate from page # to frame #?

- **We have to search the inverted page table until we find match for the page # and process ID**
- **That can be really slow....**
- **What to do?**
 - **Use hash table to limit the search to one — or at most a few — page-table entries**

Inverted Page Table Architecture



Segmentation versus Pagination

■ **Paging**

- Not visible to the user
- Suffers from internal fragmentation (be only $\frac{1}{2}$ frame per process)
- Address translation is simple (uses concatenation)


■ **Segmentation**

- Based on the user view
 - Since segment is a logical unit for protection and sharing, these functions easier with segmentation.
 - Requires more expensive material for address translation (requires addition).
- ## ■ **Fortunately segmentation and paging can be combined.**
- Divide segments into pages.




MULTICS Address Translation Scheme

Segment page entry contains pointer to the page table of that segment



Intel 30386 Address Translation

Intel 30386 has segmentation with two level hierarchical paging



Linux on Intel 80x86

- **Uses minimal segmentation to keep memory management implementation more portable**
- **Uses 6 segments:**
 - Kernel code
 - Kernel data
 - User code (shared by all user processes, using logical addresses)
 - User data (likewise shared)
 - Task-state (per-process hardware context)
 - LDT
- **Uses 2 protection levels:**
 - Kernel mode
 - User mode

